

[AI & Vectors](#) > [JavaScript Examples](#) > [Generate Embeddings](#) >

# Generate Embeddings

Generate text embeddings using Edge Functions.

This guide will walk you through how to generate high quality text embeddings in [Edge Functions](#) using its built-in AI inference API, so no external API is required.

## Build the Edge Function

Let's build an Edge Function that will accept an input string and generate an embedding for it. Edge Functions are server-side TypeScript HTTP endpoints that run on-demand closest to your users.

### 1 Set up Supabase locally

Make sure you have the latest version of the [Supabase CLI installed](#).

Initialize Supabase in the root directory of your app and start your local stack.

```
1 supabase init
2 supabase start
```



### 2 Create Edge Function

Create an Edge Function that we will use to generate embeddings. We'll call this `embed` (you can name this anything you like).

This will create a new TypeScript file called `index.ts` under `./supabase/functions/embed`.

```
1 supabase functions new embed
```

### 3 Setup Inference Session

Let's create a new inference session to be used in the lifetime of this function. Multiple requests can use the same inference session.

Currently, only the `gte-small` (<https://huggingface.co/Supabase/gte-small>) text embedding model is supported in Supabase's Edge Runtime.

```
./supabase/functions/embed/index.ts
```

```
1 const session = new Supabase.ai.Session('gte-small');
```

### 4 Implement request handler

Modify our request handler to accept an `input` string from the POST request JSON body.

Then generate the embedding by calling `session.run(input)`.

```
./supabase/functions/embed/index.ts
```

```
1 Deno.serve(async (req) => {
2   // Extract input string from
3   const { input } = await req.
4
5   // Generate the embedding fr
6   const embedding = await sess
7     mean_pool: true,
8     normalize: true,
9   });
10
11   // Return the embedding
12   return new Response(
13     JSON.stringify({ embedding
14     { headers: { 'Content-Type
15   });
16 });
```

Note the two options we pass to `session.run()` :

`mean_pool` : The first option sets `pooling` to `mean` . Pooling refers to how token-level embedding representations are compressed into a single sentence embedding that reflects the meaning of the entire sentence. Average pooling is the most common type of pooling for sentence embeddings.

`normalize` : The second option tells to normalize the embedding vector so that it can be used with distance measures like dot product. A normalized vector means its length (magnitude) is 1 - also referred to as a unit vector. A vector is normalized by dividing each element by the vector's length (magnitude), which maintains its direction but changes its length to 1.

5

## Test it!

To test the Edge Function, first start a local functions server.

Then in a new shell, create an HTTP request using cURL and pass in your input in the JSON body.

```
1 curl --request POST 'http://localhost:54321/functions/v1/embed' \  
2 --header 'Authorization: Bearer ANON_KEY' \  
3 --header 'Content-Type: application/json' \  
4 --data '{ "input": "hello world" }'
```



Be sure to replace `ANON_KEY` with your project's anonymous key. You can get this key by running `supabase status` .

## Next steps

Learn more about [embedding concepts](#)

[Store your embeddings](#) in a database

Edit this page on GitHub [↗](#)

🔗 Need some help? [Contact support](#)

🧪 Latest product updates? [See Changelog](#)

✅ Something's not right? [Check system status](#)

---

© Supabase Inc

—

Contributing

Author Styleguide

Open Source

SupaSquad

Privacy Settings

